

# assignemt01

September 26, 2025

## 1 Exercise-1

### 1.1 1.1 Data Exploration

#### 1.1.1 Q1.1.1

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	\
0	7.4	0.70	0.00	1.9	0.076	
1	7.8	0.88	0.00	2.6	0.098	
2	7.8	0.76	0.04	2.3	0.092	
3	11.2	0.28	0.56	1.9	0.075	
4	7.4	0.70	0.00	1.9	0.076	

	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	\
0	11.0	34.0	0.9978	3.51	0.56	
1	25.0	67.0	0.9968	3.20	0.68	
2	15.0	54.0	0.9970	3.26	0.65	
3	17.0	60.0	0.9980	3.16	0.58	
4	11.0	34.0	0.9978	3.51	0.56	

	alcohol	quality	Id
0	9.4	5	0
1	9.8	5	1
2	9.8	5	2
3	9.8	6	3
4	9.4	5	4

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 1143 entries, 0 to 1142
```

```
Data columns (total 13 columns):
```

#	Column	Non-Null Count	Dtype
0	fixed acidity	1143 non-null	float64
1	volatile acidity	1143 non-null	float64
2	citric acid	1143 non-null	float64
3	residual sugar	1143 non-null	float64
4	chlorides	1143 non-null	float64
5	free sulfur dioxide	1143 non-null	float64
6	total sulfur dioxide	1143 non-null	float64

7	density	1143 non-null	float64
8	pH	1143 non-null	float64
9	sulphates	1143 non-null	float64
10	alcohol	1143 non-null	float64
11	quality	1143 non-null	int64
12	Id	1143 non-null	int64

dtypes: float64(11), int64(2)

memory usage: 116.2 KB

	fixed acidity	volatile acidity	citric acid	residual sugar \
count	1143.000000	1143.000000	1143.000000	1143.000000
mean	8.311111	0.531339	0.268364	2.532152
std	1.747595	0.179633	0.196686	1.355917
min	4.600000	0.120000	0.000000	0.900000
25%	7.100000	0.392500	0.090000	1.900000
50%	7.900000	0.520000	0.250000	2.200000
75%	9.100000	0.640000	0.420000	2.600000
max	15.900000	1.580000	1.000000	15.500000

	chlorides	free sulfur dioxide	total sulfur dioxide	density \
count	1143.000000	1143.000000	1143.000000	1143.000000
mean	0.086933	15.615486	45.914698	0.996730
std	0.047267	10.250486	32.782130	0.001925
min	0.012000	1.000000	6.000000	0.990070
25%	0.070000	7.000000	21.000000	0.995570
50%	0.079000	13.000000	37.000000	0.996680
75%	0.090000	21.000000	61.000000	0.997845
max	0.611000	68.000000	289.000000	1.003690

	pH	sulphates	alcohol	quality	Id
count	1143.000000	1143.000000	1143.000000	1143.000000	1143.000000
mean	3.311015	0.657708	10.442111	5.657043	804.969379
std	0.156664	0.170399	1.082196	0.805824	463.997116
min	2.740000	0.330000	8.400000	3.000000	0.000000
25%	3.205000	0.550000	9.500000	5.000000	411.000000
50%	3.310000	0.620000	10.200000	6.000000	794.000000
75%	3.400000	0.730000	11.100000	6.000000	1209.500000
max	4.010000	2.000000	14.900000	8.000000	1597.000000

### 1.1.2 Q1.1.2

total sulfur dioxide	32.782130
free sulfur dioxide	10.250486
fixed acidity	1.747595
residual sugar	1.355917
alcohol	1.082196
quality	0.805824

```

citric acid          0.196686
volatile acidity     0.179633
sulphates            0.170399
pH                   0.156664
chlorides            0.047267
density              0.001925
Name: std, dtype: float64

```

We can see that “*total sulfur dioxide*” and “*free sulfur dioxide*” have the highest variation in the dataset. This indicates that sulfur content varies a lot between the tests. Features like *chlorides* and *density*, on the other hand, have low variation, which suggests that the values for these features are more stable.

## 1.2 Correlation Analysis

### 1.2.1 Q1.2.1

	fixed acidity	volatile acidity	citric acid	\
fixed acidity	1.000000	-0.250728	0.673157	
volatile acidity	-0.250728	1.000000	-0.544187	
citric acid	0.673157	-0.544187	1.000000	
residual sugar	0.171831	-0.005751	0.175815	
chlorides	0.107889	0.056336	0.245312	
free sulfur dioxide	-0.164831	-0.001962	-0.057589	
total sulfur dioxide	-0.110628	0.077748	0.036871	
density	0.681501	0.016512	0.375243	
pH	-0.685163	0.221492	-0.546339	
sulphates	0.174592	-0.276079	0.331232	
alcohol	-0.075055	-0.203909	0.106250	
quality	0.121970	-0.407394	0.240821	

	residual sugar	chlorides	free sulfur dioxide	\
fixed acidity	0.171831	0.107889	-0.164831	
volatile acidity	-0.005751	0.056336	-0.001962	
citric acid	0.175815	0.245312	-0.057589	
residual sugar	1.000000	0.070863	0.165339	
chlorides	0.070863	1.000000	0.015280	
free sulfur dioxide	0.165339	0.015280	1.000000	
total sulfur dioxide	0.190790	0.048163	0.661093	
density	0.380147	0.208901	-0.054150	
pH	-0.116959	-0.277759	0.072804	
sulphates	0.017475	0.374784	0.034445	
alcohol	0.058421	-0.229917	-0.047095	
quality	0.022002	-0.124085	-0.063260	


  

	total sulfur dioxide	density	pH	sulphates	\
--	----------------------	---------	----	-----------	---

fixed acidity	-0.110628	0.681501	-0.685163	0.174592
volatile acidity	0.077748	0.016512	0.221492	-0.276079
citric acid	0.036871	0.375243	-0.546339	0.331232
residual sugar	0.190790	0.380147	-0.116959	0.017475
chlorides	0.048163	0.208901	-0.277759	0.374784
free sulfur dioxide	0.661093	-0.054150	0.072804	0.034445
total sulfur dioxide	1.000000	0.050175	-0.059126	0.026894
density	0.050175	1.000000	-0.352775	0.143139
pH	-0.059126	-0.352775	1.000000	-0.185499
sulphates	0.026894	0.143139	-0.185499	1.000000
alcohol	-0.188165	-0.494727	0.225322	0.094421
quality	-0.183339	-0.175208	-0.052453	0.257710

	alcohol	quality
fixed acidity	-0.075055	0.121970
volatile acidity	-0.203909	-0.407394
citric acid	0.106250	0.240821
residual sugar	0.058421	0.022002
chlorides	-0.229917	-0.124085
free sulfur dioxide	-0.047095	-0.063260
total sulfur dioxide	-0.188165	-0.183339
density	-0.494727	-0.175208
pH	0.225322	-0.052453
sulphates	0.094421	0.257710
alcohol	1.000000	0.484866
quality	0.484866	1.000000

### 1.2.2 Q1.2.2



### 1.2.3 Q1.2.3

By looking at the plot from **Q1.2.2**, the feature with the strongest positive correlation with quality is alcohol (0.48), indicating that wines with higher alcohol content tend to have higher quality. The feature with the strongest negative correlation with quality is volatile acidity (-0.41), meaning that higher volatile acidity is associated with lower wine quality.

### 1.2.4 Q1.2.4

The correlation between pH and quality is -0.052. This indicates that there is almost no relationship between these two features, meaning that pH has very little effect on wine quality.

The correlation between alcohol and quality is 0.48, which is much higher than for pH. This suggests that wines with higher alcohol content tend to have better quality, making alcohol a better predictor of wine quality than pH.

---

## 1.3 1.3 Linear Regression

### 1.3.1 Q1.3.1

Intercept: [5.62647607]

Coefficient: [0.40830322]

```
/home/shekhe9920/ml/my_env/lib/python3.10/site-  
packages/sklearn/utils/validation.py:1406: DataConversionWarning: A column-  
vector y was passed when a 1d array was expected. Please change the shape of y  
to (n_samples, ), for example using ravel().
```

```
y = column_or_1d(y, warn=True)
```

### 1.3.2 Q1.3.2

Intercept: [1.88701286]

Coefficient: [[0.36104097]]

### 1.3.3 Q1.3.3

Model using **chlorides**: \* Intercept: 5.62647607 \* Coefficient: 0.40830322

The positive coefficient indicates that higher chloride levels are associated with slightly higher wine quality. But, this relationship is weak and may not be meaningful in practice.

Model using **alcohol**: \* Intercept: 1.88701286 \* Coefficient: 0.36104097

The positive coefficient shows that wines with higher alcohol content tend to have higher quality.

Comparison: The model using alcohol is a better predictor of wine quality because the relationship is stronger and more meaningful. While the correlation between chlorides and quality is slightly negative (-0.12), the regression model shows that chlorides are not a reliable predictor of wine quality.

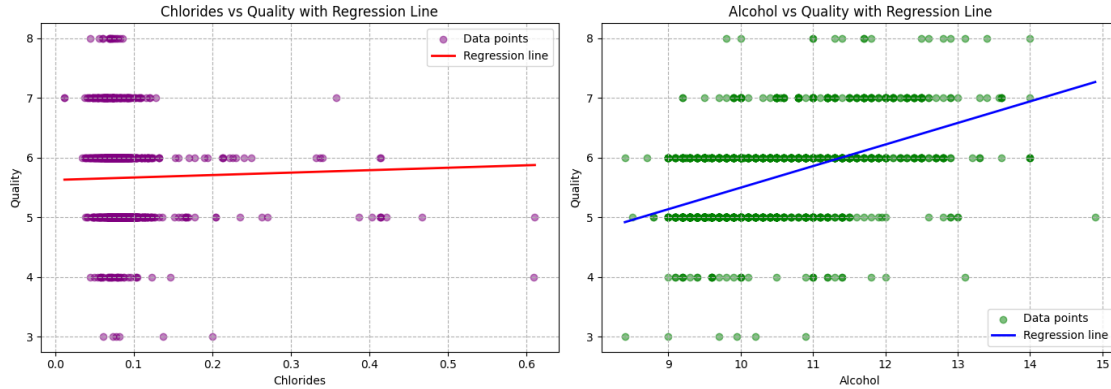
### 1.3.4 Q1.3.4

```
/home/shekhe9920/ml/my_env/lib/python3.10/site-  
packages/sklearn/utils/validation.py:2749: UserWarning: X does not have valid  
feature names, but SGDRegressor was fitted with feature names
```

```
warnings.warn(
```

```
/home/shekhe9920/ml/my_env/lib/python3.10/site-  
packages/sklearn/utils/validation.py:2749: UserWarning: X does not have valid  
feature names, but LinearRegression was fitted with feature names
```

```
warnings.warn(
```



**Chlorides:** The line does not match the data well. The correlation between quality and chlorides is too weak, so the model gives almost the same prediction for all values of chlorides.

**Alcohol:** The line fits the data much better. It shows a clear upward trend: the higher the alcohol content, the better the quality. This agrees with the correlation of 0.48.

## 1.4 1.4 Train-Test Split

### 1.4.1 Q1.4.1

Fold 1

MSE: 0.4241, RMSE: 0.6513,  $R^2$ : 0.3048

-----

Fold 2

MSE: 0.4592, RMSE: 0.6777,  $R^2$ : 0.2680

-----

Fold 3

MSE: 0.4924, RMSE: 0.7017,  $R^2$ : 0.2635

-----

Fold 4

MSE: 0.5253, RMSE: 0.7248,  $R^2$ : 0.2928

-----

Fold 5

MSE: 0.4513, RMSE: 0.6718,  $R^2$ : 0.2965

-----

Average  $R^2$  across 5 splits: 0.2851

By looking at the  $R^2$  score (explained variance), we can see how well alcohol alone can predict wine quality. A higher  $R^2$  score means the model has a better chance of predicting the quality.

- $R^2$  **close to 1** → the model predicts very well.
- $R^2$  **close to 0** → the model does not predict well.

From the results above, we got an  $R^2$  score of about **0.285**. This means alcohol alone explains about **28.5% of the variation** in wine quality.

This shows that alcohol has some relationship with wine quality, but it is **not enough on its own**. Wine quality also depends on many other factors.

---

### 1.4.2 Q1.4.2

Fold 1

MSE: 0.5918, RMSE: 0.7693,  $R^2$ : 0.0299

-----

Fold 2

MSE: 0.6297, RMSE: 0.7936,  $R^2$ : -0.0038

-----

Fold 3

MSE: 0.6620, RMSE: 0.8137,  $R^2$ : 0.0098

-----

Fold 4

MSE: 0.7264, RMSE: 0.8523,  $R^2$ : 0.0220

-----

Fold 5

MSE: 0.6394, RMSE: 0.7996,  $R^2$ : 0.0034

-----

Average  $R^2$  across 5 splits: 0.0123

The  $R^2$  score is much lower this time when we only use **chlorides** to predict wine quality. It is almost **0 (0.0123)**, which means that the model does not predict well. This shows that chlorides alone have almost **no relationship** with wine quality, and many other factors are needed to make good predictions.

### 1.4.3 Q1.4.3

Yes, the model is underfitting in both cases mentioned in question Q1.4.1 and Q1.4.2. This is because using only **alcohol** or **chlorides** (especially chlorides) is not enough to make accurate predictions for wine quality. With so little information, the model becomes **too simple** and produces **unreasonable outputs**, since it cannot capture the true complexity of the data.

### 1.4.4 Q1.4.4

Alcohol:

Mean  $R^2$ : 0.2851

Variance of  $R^2$ : 0.0003

Chloride:

Mean  $R^2$ : 0.0123

Variance of  $R^2$ : 0.0001

The results show that **alcohol** has a much higher mean  $R^2$  (0.2851) compared to **chlorides** (0.0123). This means alcohol alone explains about **28.5% of the variation** in wine quality, while chlorides explain almost nothing.

The variance for alcohol (0.0003) is slightly higher than for chlorides (0.0001), which shows that the performance with alcohol varies a little between folds, while chlorides consistently give very



poor results.

In conclusion, alcohol is a much stronger predictor of wine quality than chlorides, but even alcohol alone is not enough to create an accurate model. Wine quality depends on multiple factors, not just one feature.

---

## 1.5 1.5 Multiple Linear Regression

### 1.5.1 Q1.5.1

```
Fold 1
MSE: 0.3574, RMSE: 0.5979, R2: 0.4141
-----
Fold 2
MSE: 0.3646, RMSE: 0.6038, R2: 0.4188
-----
Fold 3
MSE: 0.4138, RMSE: 0.6433, R2: 0.3811
-----
Fold 4
MSE: 0.4450, RMSE: 0.6671, R2: 0.4009
-----
Fold 5
MSE: 0.3911, RMSE: 0.6254, R2: 0.3903
-----
```

### 1.5.2 Q1.5.2

From the output results of question Q1.5.1, we can see that the model performs better in each fold compared to the models that only used **alcohol** or **chlorides** as a single feature to predict wine quality.

The **R<sup>2</sup> score** is much higher, which indicates a more accurate prediction of wine quality.

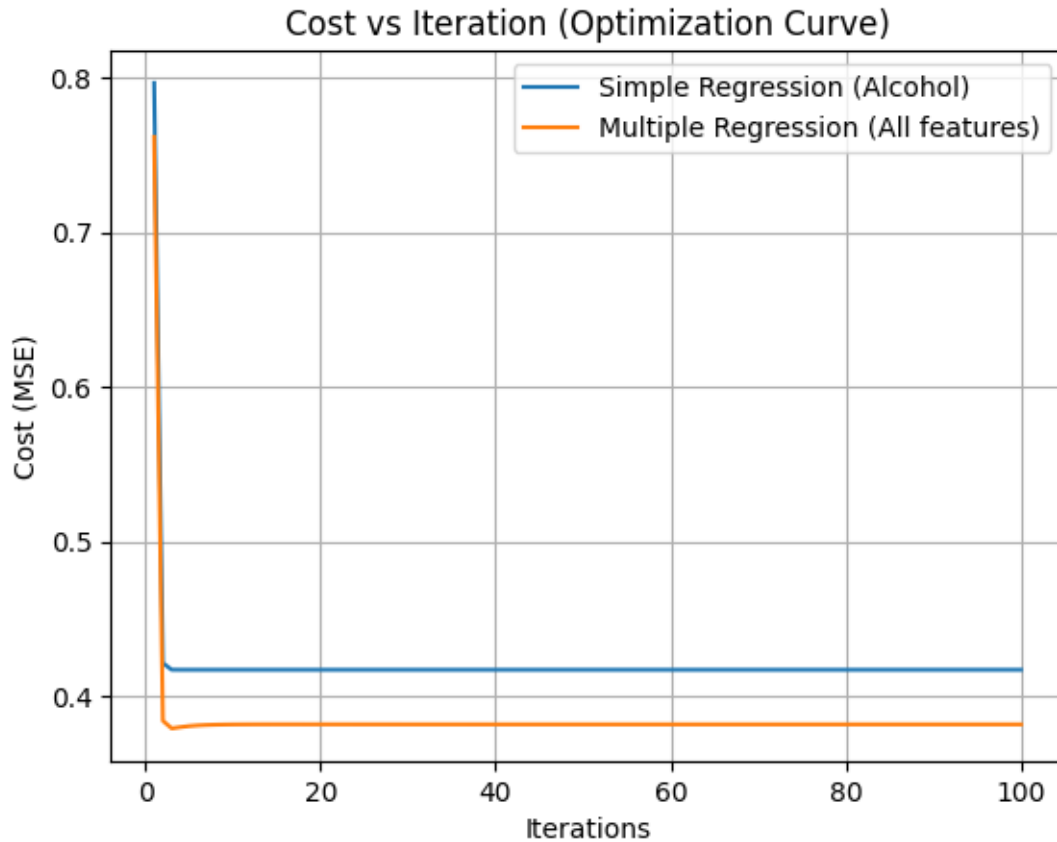
We can also see that the **MSE** and **RMSE** values are lower, which is a good sign since it means the model is making fewer errors in its predictions.

Here is the calculation of the mean and variance of the R<sup>2</sup> score for the multiple linear regression model:

```
Mean R2: 0.4011
Variance of R2: 0.0002
```

As we can see, the mean  $R^2$  score is much higher than both models from question Q1.4.4. The variance of the  $R^2$  score is between the variance of the models using alcohol and chlorides, showing that this model is more stable. By looking at both the mean and variance, we can conclude that this multiple linear regression model is much more stable and much better at predicting wine quality than using just one feature.

### 1.5.3 Q1.5.3



### 1.5.4 Q1.5.4

The *multiple linear regression* performs better.

*Why:* - It achieves a higher  $R^2$  ( $\approx 0.40$ ) than the simple models

- Alcohol only:  $R^2 \approx 0.29$

- Chlorides only:  $R^2 \approx 0.01 \rightarrow$  Higher  $R^2$  means it explains more variation in wine quality. - It has lower MSE/RMSE (smaller prediction errors). - In the **Cost vs Iteration** plot, it converges to a lower final cost, showing a better fit. - Its variance across folds is low, indicating stable performance.

Since wine quality depends on multiple variables, using all features provides a more accurate and stable model than using a single feature.


---

---

## 2 Exercise-2

### 2.0.1 Q2.1

	Feature	Importance
10	alcohol	0.282929
9	sulphates	0.134164
1	volatile acidity	0.133508
6	total sulfur dioxide	0.070300
8	pH	0.059995
4	chlorides	0.058002
0	fixed acidity	0.057861
2	citric acid	0.054573
7	density	0.051610
3	residual sugar	0.049845
5	free sulfur dioxide	0.047214



As we can see from the plot above, **alcohol** is the most influential feature in predicting wine quality. This is a reasonable result because, in Q1.2.2, we observed that alcohol had the highest correlation with quality among all the features.

### 2.0.2 Q2.2

#### a.) Polynomial regression:

	Degree	Train RMSE	Test RMSE	R <sup>2</sup> (Test)
0	1	0.644291	0.616468	0.317069
1	2	0.602874	0.632594	0.280872

2	3	0.441140	1.927599	-5.677121
3	4	0.045866	90.644456	-14764.168534

#### Degree 1 → Degree 2

- **Test RMSE** increased slightly ( $0.6165 \rightarrow 0.633$ ) → performance **worsened**, not improved.
- $R^2$  decreased ( $0.31707 \rightarrow 0.281$ ) → the model **explained less variance** in wine quality.

*Conclusion:* Adding quadratic terms did **not improve performance**. It slightly **worsened generalization**.

#### Degrees 3 and 4

- **Train RMSE** dropped massively, reaching almost **0** at degree 4 → the model **memorized the training set**.
- **Test RMSE** skyrocketed (up to **90+**) and  $R^2$  became extremely negative → **overfitting**.

*Conclusion:* Higher polynomial degrees **overfit** the data, making the model **unusable for unseen data**.

#### b.) Regularization:

	Model	Alpha	Train RMSE	Test RMSE	$R^2$ (Test)
0	Ridge	0.0001	0.644302	0.616150	0.317774
1	Ridge	0.0010	0.644466	0.615400	0.319433
2	Ridge	0.0100	0.644640	0.615050	0.320206
3	Ridge	0.1000	0.644691	0.614499	0.321424
4	Ridge	1.0000	0.645511	0.611074	0.328968
5	Lasso	0.0001	0.644682	0.614584	0.321236
6	Lasso	0.0010	0.645466	0.611158	0.328782
7	Lasso	0.0100	0.655957	0.603731	0.344998
8	Lasso	0.1000	0.714413	0.633782	0.278168
9	Lasso	1.0000	0.807366	0.729726	0.043079

Coefficient comparison (Ridge vs Lasso):

	Feature	Ridge (alpha=0.05)	Lasso (alpha=0.05)
9	sulphates	0.919276	0.000000
10	alcohol	0.296000	0.315921
0	fixed acidity	0.023085	0.047236
5	free sulfur dioxide	0.002125	0.005235
6	total sulfur dioxide	-0.002348	-0.003363
3	residual sugar	-0.009001	-0.000000
7	density	-0.296994	-0.000000
2	citric acid	-0.344764	0.000000
8	pH	-0.398887	-0.000000
1	volatile acidity	-1.362192	-0.000000
4	chlorides	-1.789333	-0.000000

#### 1. Generalization (Train/Test RMSE and ( $R^2$ )):

##### Ridge:

- As alpha increases, the Test RMSE slowly goes down (from  $0.616 \rightarrow 0.611$ ), and  $R^2$  goes up (from  $0.318 \rightarrow 0.329$ ).

- This means Ridge **reduces overfitting** and improves how well the model works on new data.
- Ridge never gets much worse, even at higher alpha values.

#### Lasso:

- At low alpha ( $0.0001 \rightarrow 0.01$ ), Lasso **improves** Test RMSE and  $R^2$ , reaching its best  $R^2 = 0.345$ .
- When alpha gets too high ( $0.1 \rightarrow 1.0$ ), performance **drops a lot**. Test RMSE increases, and  $R^2$  goes close to **0**.
- This happens because Lasso becomes too strict and **removes too many features**, causing underfitting.

#### 2. Coefficients: Ridge:

- Shrinks all coefficients smoothly but keeps them non-zero.
- *Example:* even small features like “free sulfur dioxide” still have a small effect.

#### Lasso:

- Forces many coefficients to exactly zero, leaving only the most important features like alcohol.
- *Example:* sulphates, citric acid, density, and others are completely removed (set to 0).

#### c.) Model comparison:

Best parameters for Random Forest: {'max\_depth': None, 'min\_samples\_leaf': 1, 'min\_samples\_split': 2, 'n\_estimators': 200}

#### Random Forest Performance:

	Model	Train RMSE	Test RMSE	$R^2$ (Test)
0	Random Forest	0.226678	0.548983	0.458405

#### 2.0.3 Model Comparison

**Linear Regression:** -  $R^2 \approx 0.40$  - Test  $RMSE \approx 0.61$  - Assumes a straight-line relationship between the features and wine quality.

**Random Forest:** -  $R^2 \approx 0.46$  - Test  $RMSE \approx 0.55$  - Can capture complex patterns and interactions between features.

#### 2.0.4 Which performs better and why?

The **Random Forest** model performs **better** because:

- It has **lower Test RMSE** and **higher  $R^2$**   $\rightarrow$  more accurate predictions.
- It does **not assume a linear relationship**, so it can model more complex patterns in the data.
- By combining many decision trees, it becomes **more robust** and handles noise better.

*Conclusion:* Random Forest is a better choice for this dataset because wine quality depends on many factors that interact in non-linear ways.